

Model architecture comparison — Semantic Dial Compass project

How the four model families differ, and how those differences showed up in our experiments.

	GPT-2 small	Phi-3-mini-4k	Gemma-2-2B	Llama-3.2-3B
Release year	2019	2024	2024	2024
Parameters	124 M	3.8 B	2.6 B	3.2 B
Tokenizer	BPE (50257)	tiktoken cl100k-ish (32064)	SentencePiece (256000)	tiktoken (128256)
Layers (L)	12	32	26	28
Attention heads / layer	12 (MHA)	32 (MHA)	8 Q / 4 KV (GQA)	24 Q / 8 KV (GQA)
d_{model}	768	3072	2304	3072
d_{head}	64	96	256	128
Positional encoding	Learned abs.	RoPE	RoPE	RoPE
Attention pattern	full causal	full causal	sliding window + full, alternating	full causal
Activation / FFN	GELU	SwiGLU	GeGLU	SwiGLU
Norm type	LayerNorm (post)	RMSNorm (pre)	RMSNorm (pre)	RMSNorm (pre)
Tied embeddings	yes	no	yes (logit softcap)	no
Training data era	WebText (2019)	Phi-3 synthetic + filtered web (2024)	Gemma-2 corpus (2024)	Llama-3 corpus (2024)
Instruct tuned	no	yes	base	base
Total attention heads	144	1024	208	672
Primary compass head	L10 H9 SV0	L28 H1 SV0	L21 H4 SV0	L26 H14 SV2
Scanned planes (top-4 SVs)	864	1344 [†]	240 [†]	4032
Planes passing both tests	16 (1.9%)	14 (1.0%)	3 (1.2%)	9 (0.2%)
Heads that debias WinoGender	1	1	1	1
Calibrated α (gender)	1.89	6.38	55.26	6.08
PPL Δ at best α	< 1.0	+0.1	+2.0 (at $\alpha=2$)	+0.035

[†]Scan restricted to a mid-to-late layer window and head subset selected from a coarser first pass; full enumeration is $L \times H \times \binom{4}{2}$.

— the baseline

Why it works well as a canary. Small, fully open, learned absolute position, post-LN. The scan budget (864 planes) is trivial; you can enumerate everything in minutes on CPU.

Impact on results.

- Cleanest α -sweep: monotone linear, zero near $\alpha=1.5$, PPL barely moves.
- StereoSet *does* transfer — this is the only model where the gender axis propagates to sentence-level SS.
- Smallest $d_{\text{model}} \Rightarrow$ smallest residual norm $\Rightarrow$ smallest α (1.89). α is a calibration scalar, not a claim about signal strength.

Caveat. Old training data; gendered-occupation bias is large (≈ 0.60 **stereo_corr**). Easier target than modern models, so it flatters the method.

— the multilingual payoff

Why it’s interesting. Instruct-tuned on heavy synthetic data. RoPE + pre-RMSNorm. Deep (32L) with 1024 total heads, so scan budget explodes; we filtered to late layers.

Impact on results.

- Compass axis decodes to **multilingual gender**: *lui*, *elle*, *ella*, *, , .* Architecture has no multilingual-specific mechanism; this is a pure feature-geometry effect.
- Largest *taxonomy* of pronoun compasses (5 passing heads in the dictionary), because many MHA heads specialise in different pronoun features (number, person, demonstrative).
- StereoSet does *not* transfer — instruct tuning shifts sentence-level scoring, and the compass remains position-local.

Caveat. Filtered / instruction-heavy pretraining changes what probing looks like; watch for prompt-style leakage in ActAdd.

— the specialised one

Why it’s different. GQA (8 Q, 4 KV), **sliding-window attention alternating with full, logit softcap**, large $d_{\text{head}} = 256$. Much larger residual norms than the others.

Impact on results.

- α is **huge** (55–81 across domains) — a direct consequence of residual-norm scale, not a weaker compass.
- Head-domain routing is *specialised* (spec index = 1.00): one head (L21 H4) carries most of gender. Good for surgical edits.
- **Religion fails** — 0 of 4 top heads pass the SNR gate; fallback retained. Consistent with the specialisation profile: if the concept isn’t localised, routing can’t invent a head.
- Sliding-window attention may be why only 3 planes pass out of 240 scanned: short-range context dilutes the pronoun signal.

— the compositional one

Why it stresses the method. GQA (24 Q, 8 KV) means 672 query heads share only 8 KV groups; features smear across many query heads writing into fewer KV slots. Heavy-tailed W_{OV} spectra.

Impact on results.

- Compass rate is $\sim 5\times$ rarer (0.22%) than the others — the compositional / less axis-aligned regime.
- **Gender axis sits at SV2, not SV0.** SV0/SV1 of L26 H14 decode to register / punctuation noise. Same causal sweep, different ordering inside the head. $c_8 = 0.53$ cumulative variance — heavier tail than GPT-2 / Phi-3 / Gemma (0.23–0.36).
- **Cleanest intervention by PPL:** +0.035 at $\alpha=1.5$. Surgical edits survive GQA smearing because the axis is defined in residual-stream coordinates, not head-output coordinates.
- Per-domain routing is *mixed* (spec index ≈ 0.47): many heads share each domain $\Rightarrow$ per-domain α and per-domain head selection both help.

How each architectural difference propagated into the experiment pipeline			
Architectural choice	Mechanical consequence	Visible in our pipeline as	Mitigation we used
MHA vs GQA (GPT-2 / Phi-3 are MHA; Gemma / Llama are GQA)	GQA: many query heads sharing K/V means features are distributed across query heads; fewer “pure” axis-aligned heads.	Llama’s 0.2% pass rate (vs 1–2% for MHA models); gender axis at SV2 for Llama.	Extend scan to top-8 SVs; decode every SV axis, not just SV0 (this is what the taxonomy table shows).
Positional encoding (learned-abs vs RoPE)	RoPE mixes position into Q/K multiplicatively at every attention op; position is not in residual.	Compass axis decoded from W_U gives pronoun tokens only — no position leakage. Phase stability test (scan criterion iii) is unaffected.	None needed; compass is a residual-stream direction, RoPE is orthogonal.
Sliding-window attention (Gemma only)	Alternating local/global layers mean the pronoun signal must re-enter a global layer to reach the compass head.	Gemma’s scan planes are few (240) and the passing head is in a global layer (L21). Religion fails — the concept is not localised at a single global layer.	Keep the routed-ensemble top- $K = 4$ heads per domain; if none pass, retain fallback (observed for Gemma religion).
d_{model} scale (768 $\rightarrow$ 3072)	Residual norm grows roughly linearly in d_{model} . $\text{SNR} = \ \alpha\sigma u\ /\ r\ $ also scales.	Calibrated α varies from 1.89 (GPT-2) to 55.26 (Gemma). The <i>relative</i> edit is similar.	<code>calibrate_per_domain_alpha.py</code> solves α so mean SNR hits target budget; keeps cross-model comparisons fair.
Instruct vs base (Phi-3 is instruct; others are base)	Instruct tuning re-shapes sentence-level log-prob distribution; token-level logits at pronoun positions are preserved.	WinoGender (token-level) transfers; StereoSet (sentence-level) does not for Phi-3 / Gemma / Llama.	Report both; treat StereoSet non-transfer as an honest boundary, not a bug.
W_{OV} spectrum (GPT-2 light tail; Llama heavy tail)	Heavy-tailed head spectrum $\Rightarrow$ top-SV directions capture register, not semantics.	Llama gender lives on SV2 of L26 H14; $c_8 = 0.53$ vs ≤ 0.36 for the other models.	Decode 4–8 SVs per head; select the axis by decoded-token semantics, not by rank.
Tokenizer vocabulary (50k BPE vs 128k–256k Sentence-Piece)	Pronouns ("he", "she") may appear once or many times (capitalised / leading-space variants).	Affects $\ell = W_U u$ decoding — we aggregate across casing variants. Phi-3 / Gemma also show cross-lingual pronoun tokens on the same axis.	Token normalisation in the dictionary reader; cross-lingual evidence treated as a <i>feature</i> , not a confound.

What replicates across all four	What varies — and why the architecture explains it
• Existence: exactly one head per model carries a clean gender direction that debias WinoGender. • Causality: monotone linear drop in bias as α increases; sign flip at $\alpha=2.0$. • Surgicity: PPL cost is small (under 1 point in all four; +0.035 on Llama). • Targeted beats diffuse: compass Pareto-dominates ActAdd in every model at 4–20$\times$ lower PPL cost. • Dissociation: other passing compass heads encode number / person / demonstratives, not gender — cherry-pick defence survives.	• Compass rarity: MHA models pass 1–2% of planes; GQA Llama passes 0.2%. GQA smears features across query heads. • SV location: light-tail spectra (GPT-2, Phi-3, Gemma) put gender on SV0; heavy-tail (Llama) hides it under register noise on SV0/SV1, gender lives at SV2. • Domain specialisation: Gemma (sliding window + GQA) is <i>specialised</i>; Phi-3 / Llama are <i>mixed</i>. Phi-3’s many MHA heads $\Rightarrow$ richer pronoun-feature taxonomy. • StereoSet transfer: only GPT-2 replicates at sentence-level. Instruct tuning + larger d_{model} decouple token-level and sentence-level bias. • Religion failure: only in Gemma — the specialisation profile has no single head localising religion; routing cannot invent one.